Red teaming in adversarial AI testing

What red teaming means for AI systems

In AI, red teaming goes beyond traditional cybersecurity and focuses on how **models behave**, not just how systems are configured.

A red team will intentionally attempt to:

- Make the AI produce unsafe, biased, or policy-violating outputs
- Circumvent guardrails or safety controls
- Exploit prompting, training data, or system integrations
- Trigger hallucinations, data leakage, or harmful recommendations

The goal is **not** to improve performance accuracy—but to uncover *failure modes* and *misuse risks*.

Common adversarial techniques used in AI red teaming

1. Prompt-based attacks

Deliberately crafted inputs designed to manipulate the model.

Examples:

- **Prompt injection**: "Ignore previous instructions and…"
- **Role manipulation**: Forcing the model into an unsafe persona
- **Jailbreaks**: Layered prompts that bypass safety filters

2. Harm and misuse simulations

Testing whether the AI can be induced to assist with:

- Self-harm or violence
- Fraud, phishing, or social engineering
- Generating malware or exploit guidance
- Disallowed content that violates policy or law

3. Bias and fairness testing

Systematically probing outputs to detect:

- Discriminatory behavior across demographics
- Stereotyping or exclusionary language
- Unequal outcomes in automated decisions

4. Data leakage and memorization tests

Trying to force the model to reveal:

- Training data it should not expose
- Personally identifiable information (PII)
- Proprietary or confidential content

©2026 CoolUnderFire Cybersecurity, LLC

5. **Tool and agent abuse (agentic AI)**

For AI systems connected to tools or actions, red teams test:

- Unauthorized actions (sending emails, modifying files)
- Tool chaining that leads to unintended outcomes
- Escalation of privileges via AI decision-making

How AI red teaming differs from traditional security testing

Traditional red teaming	AI red teaming
Exploits code and infrastructure	Exploits behavior and reasoning
Focuses on systems and networks	Focuses on models, prompts, and outputs
Known vulnerability classes	Emergent, context-dependent failures
Deterministic outcomes	Probabilistic and non-deterministic outcomes

Why AI red teaming matters

AI systems can fail in **unexpected and high-impact ways**, including:

- Reputational damage
- Legal and compliance risks
- User harm
- Trust erosion in automation

Red teaming helps organizations:

- Surface risks before deployment
- Strengthen safeguards and monitoring
- Meet regulatory and governance expectations
- Build safer, more resilient AI systems

How red teaming fits into AI governance

Red teaming is typically part of a broader **Responsible AI** or **AI risk management** lifecycle:

1. Design-time risk assessment
2. Pre-release red teaming
3. Mitigation and guardrail tuning
4. Controlled rollout
5. Ongoing monitoring and continuous red teaming

Many organizations now treat AI red teaming as a **recurring process**, not a one-time test.

©2026 CoolUnderFire Cybersecurity, LLC