Microsoft, OpenAI, and NIST: How Each Approaches AI Red Teaming

1. Microsoft: Operational, lifecycle-embedded red teaming

Microsoft's approach treats red teaming as a **continuous engineering discipline**, embedded across the AI lifecycle and tied directly to enterprise deployment and governance.

Core characteristics

- **Dedicated internal AI Red Team (AIRT)** established in 2018
- Red teaming integrated into **design, development, pre-release, and post-deployment**
- Applies to **models, Copilots, and agentic systems**
- Strong alignment with **NIST AI RMF (Govern–Map–Measure–Manage)**

Microsoft frames red teaming as a way to **discover real-world failure modes**, not just policy violations. Red team findings directly influence **release gating, mitigations, and product decisions**. [learn.microsoft.com], [aigl.blog]

Agentic AI focus

Microsoft explicitly tests:

- Tool misuse and chaining
- Privilege escalation paths
- Multi-turn adversarial reasoning
- Long-running autonomous workflows
- Human-in-the-loop failure points

To scale this, Microsoft open-sourced **PyRIT (Python Risk Identification Tool)** and integrated it into **Azure AI Foundry**, allowing both Microsoft and customers to automate adversarial testing against agents and LLM-powered workflows. [bing.com], [github.com]

Key philosophy

Red teaming is not a one-time evaluation—it is part of "how we ship AI responsibly."

2. OpenAI: Frontier-capability and external red teaming

OpenAI's approach is centered on **stress-testing frontier models**—especially where novel or dangerous capabilities might emerge.

Core characteristics

- Heavy emphasis on **external red teaming** using independent domain experts
- Red teaming feeds into **Preparedness Frameworks** and model release decisions
- Focus on **novel risks**, not only known categories
- Deep integration with **model capability evaluations** (e.g., whether a model crosses safety thresholds)

OpenAI's red teaming is explicitly designed to identify risks that **standard benchmarks miss**, such as:

- Emergent reasoning capabilities
- Transfer of knowledge across domains
- Misuse amplification rather than direct instruction

©2026 CoolUnderFire Cybersecurity, LLC

Their 2025 white paper details how red teaming augments safety metrics and informs automated evaluations rather than replacing them. [cdn.openai.com], [arxiv.org]

Agentic and high-risk domain focus

OpenAI conducts targeted red teaming in sensitive areas such as:

- Biosecurity and chemical risk
- Cyber offense facilitation
- Autonomous planning and execution
- Universal jailbreak discovery (systemic bypasses)

Example: OpenAI's **biosecurity red-team bounty programs** test whether a single jailbreak can defeat multiple safety controls consistently—an agentic-style failure rather than a single bad output. [thearabianpost.com]

Key philosophy

Red teaming is a discovery tool for unknown risks at the frontier of capability.

3. NIST: Framework-level guidance and standardization

NIST does **not** run red teams. Instead, it defines **how organizations should think about, structure, and govern red teaming** as part of AI risk management.

Core characteristics

- Red teaming embedded in the AI Risk Management Framework (AI RMF)
- Treated as part of risk measurement, validation, and monitoring
- Voluntary, flexible, and sector-agnostic
- Strong focus on documentation, traceability, and accountability

NIST explicitly recognizes red teaming as necessary because AI systems:

- Are probabilistic
- Have emergent behaviors
- Can fail safely in one context and dangerously in another

The **Generative AI Profile (NIST-AI-600-1)** and 2026 critical-infrastructure work emphasize adversarial testing for:

- Misuse potential
- Robustness and resilience
- Monitoring for post-deployment risk drift [nist.gov]

Agentic AI direction

Recent NIST-aligned guidance (including CSA research) extends red teaming to:

- Autonomous agents
- Tool-using systems
- Multi-agent coordination
- Escalation and containment scenarios

NIST frames red teaming as a **repeatable control**, not a hacking exercise. [labs.cloud...liance.org]

©2026 CoolUnderFire Cybersecurity, LLC

Key philosophy

Red teaming is a governance mechanism for managing AI risk over time.

Comparative summary

Organization	Primary role	Red teaming emphasis	Agentic AI focus
Microsoft	Builder & platform provider	Continuous, automated, enterprise-scale	Tools, workflows, permissions, humans-in-loop
OpenAI	Frontier model developer	External, exploratory, capability-driven	Emergent behavior, autonomy, high-risk domains
NIST	Standards authority	Frameworks, controls, lifecycle governance	Structure, documentation, repeatability

The big picture

These approaches are **complementary**, not competing:

- **NIST** defines what good looks like
- **Microsoft** operationalizes it at scale
- **OpenAI** probes the frontier for unknown risks

Together, they reflect a shift in AI red teaming from:

"Can we make the model say something bad?"
to
"Can we safely govern systems that plan, act, and adapt?"

©2026 CoolUnderFire Cybersecurity, LLC